

THONG PHAM

Data Engineer | Databricks | AWS | Spark | Kafka

206-209-9482 | phamhoanghatthong1999@gmail.com | [linkedin.com/in/pham-thong](https://www.linkedin.com/in/pham-thong)

Data Engineer with 5+ years of experience building scalable batch and streaming platforms using Databricks, AWS, Spark, Kafka, Snowflake, and Airflow. Experienced in lakehouse migration, config-driven ingestion, data quality, observability, and CI/CD, with measurable improvements in performance, cost, and reliability.

PROFESSIONAL EXPERIENCE

Transcarent, Inc.

Jan 2026 – Present

Data Engineer II

Seattle, WA

- Migrated fragmented AWS workloads supporting **60 source services** to a unified Databricks Lakehouse, implementing Medallion architecture and incremental ingestion to improve maintainability.
- Re-engineered Braze processing through query optimization and reconciliation controls, **reducing runtime from 12 hours to 15 minutes** and **mismatch rate from 50% to 0.1%**, while **saving \$5K annually** in compute.
- Replaced Fivetran with native Databricks REST API ingestion, **cutting load time from 4 hours to 5 minutes** and **eliminating \$45K** in annual licensing expense.
- Built a GitHub-to-CloudBees-to-Databricks CI/CD workflow and implemented Monte Carlo monitoring for data volume, reconciliation, freshness, and cost, reducing manual oversight by 80%.

FullThrottle.ai (AdTech)

Jul 2023 – Dec 2025

Data Engineer

Philadelphia, PA

- Developed **30+ ingestion pipelines** across S3, Lambda, ECS, API Gateway, and CloudWatch to **process 100+ GB daily** into Snowflake, PostgreSQL, and a data lake.
- Orchestrated **70+ Airflow DAGs** and optimized dependencies and execution logic, reducing runtime by 25%; added schema-evolution handling and alerts for breaking event changes.
- Designed dimensional warehouse models using star and snowflake schemas, improving query performance by 20% and enabling analytics access for four teams.
- Built an AWS backend platform with API Gateway, containerized FastAPI services on ECS, SQS, PostgreSQL, and ElastiCache, supporting scalable production services with 99.9% uptime.
- Integrated CloudWatch with New Relic dashboards and API-based error filters, establishing unified observability and **reducing troubleshooting by 3 hours per week**.

ZaloPay (FinTech)

Mar 2021 – Mar 2023

Data Engineer

Vietnam

- Built Apache NiFi and PySpark pipelines processing **500M+ records** daily from Kafka, APIs, CSV, and Parquet, delivering near-real-time data for fraud detection.
- Developed consistency validation across 70+ log types, **detecting 8% data mismatches** and missing records while improving reliability across source systems.
- Implemented DataHub for metadata discovery, lineage, and governance across the organization's data ecosystem.
- Optimized MySQL workloads and Jenkins CI/CD, reducing server cost by 10% and deployment time by 10%; **automated 100+ JMeter performance tests** for six engineering teams.

EDUCATION

Master of Science in Computer Science, Arizona State University (GPA: 3.7)

Aug 2023 – May 2025

TECHNICAL SKILLS

Cloud & Platforms: AWS (S3, Lambda, EMR, ECS, EC2, API Gateway, Redshift, Step Functions), Databricks, Snowflake, BigQuery

Data Engineering: Spark, Kafka, Airflow, dbt, NiFi, Delta Lake, Data Modeling, Data Warehousing, Data Lakes

Programming & Databases: Python, Scala, SQL, Bash, PostgreSQL, MySQL, Redis, DynamoDB

DevOps & Observability: Docker, Git, Jenkins, CloudBees, Monte Carlo, New Relic, Grafana, JMeter, CI/CD